

Miaomiao Dai

Philadelphia, PA · dmmsjtu@engineering.upenn.edu · Personal website: dmsjtum.github.io

EDUCATION

University of Pennsylvania

M.S.E. in Robotics

Aug 2026 – present

Philadelphia, PA

University of Michigan

B.S.E. in Computer Science, Summa Cum Laude (SJTU–UMich dual-degree sequence) · GPA 3.8/4.0

Aug 2024 – May 2026

Ann Arbor, MI

- Coursework: Robot Learning (A+), Algorithmic Robotics (A), Machine Learning (A), Computer Vision (A), Data Structures & Algorithms (A)

Shanghai Jiao Tong University

B.E. in Mechanical Engineering, Global College · GPA 3.6/4.0

Aug 2022 – Aug 2024, May 2026 – Aug 2026

Shanghai, China

PUBLICATIONS

- [1] **Miaomiao Dai**, Zhongqiang Ren. “HOP: Fast Differential Dynamic Programming for Horizon-Optimal Trajectory Planning.” *Robotics: Science and Systems (RSS)*, 2026.
- [2] **Miaomiao Dai**, Zhongqiang Ren. “AL-HOP: Augmented-Lagrangian Horizon-Optimal Planning for Fast Constrained Trajectory Optimization.” In submission to *IEEE International Conference on Robotics and Automation (ICRA)*, 2027.
- [3] PhysLA: physical supervision for latent-action VLA pretraining (full title and author list withheld during double-blind review). Under review at *ICLR 2027*; **second author**.

RESEARCH EXPERIENCE

HOP: Horizon-Optimal Trajectory Planning

Jul 2025 – Feb 2026

Advisor: Prof. Zhongqiang Ren, Shanghai Jiao Tong University · RSS 2026 talk, undergraduate first author

- Proposed **HOP**, a new algorithm that makes LQR and DDP **horizon-optimal**: it optimizes the trajectory and its horizon together instead of fixing the horizon in advance.
- Proved the Riccati recursion is a **linear fractional transformation**, so all candidate horizons share one computation, $O(N^2) \rightarrow O(N)$ even when time-varying; HOP-DDP handles nonlinear systems, up to **40×** faster than brute force.
- Designed and iterated the experimental evaluation across cartpole, Segway, and quadrotor systems.

AL-HOP: Constrained Horizon-Optimal Planning

Mar 2026 – present

Advisor: Prof. Zhongqiang Ren, Shanghai Jiao Tong University · first-author manuscript in submission to ICRA 2027

- Extends HOP to hard state–control constraints (obstacles, actuator limits) with an augmented-Lagrangian formulation that keeps the one-sweep horizon prediction.
- Produced solutions valid under the full nonlinear dynamics on **100/100** benchmark problems across 10 dynamical systems, vs. 79, 71, and 55 for ALTRO and direct collocation (IPOPT, SNOPT); where both were valid, baselines needed $1.6\text{--}3.1\times$ the solve time.

PhysLA: Transferring Physical Supervision to Latent-Action VLA Pretraining

May 2026 – Sep 2026

Researcher, Shanghai Jiao Tong University · second author, under review at ICLR 2027

- Latent-action pretraining supervises contact and force only indirectly, and tactile or torque sensing covers only a small slice of video. PhysLA learns physical targets where sensors exist, infers them from causal RGB elsewhere, and uses them with visual prediction to supervise latent actions — **no tactile or torque sensing at deployment**.
- **Designed two of its three stages and co-implemented them**: learning horizon-structured physical targets from pressure, tactile, and torque measurements, and transferring them to sensor-free video via causal RGB.
- Outperforms VLA-JEPA, the strongest latent-action baseline tested, on LIBERO and every SimplerEnv task group, and beats both VLA-JEPA and $\pi_{0.5}$ on each of five real-world manipulation tasks.

Voice-Interactive Closed-Loop VLA Manipulation (B.E. thesis)

Feb 2026 – Aug 2026

Advisors: Prof. Hesheng Wang & Prof. Yunlong Guo, Shanghai Jiao Tong University · **group leader**, team of 5

- **Collected and trained on real robot data**: teleoperated **316 episodes** across 11 tabletop tasks on a low-cost SO-101, then fine-tuned $\pi_{0.5}$ on them — LoRA on the 2B vision–language backbone, full training of the 300M action expert.
- **Closed the execution loop**: a VLA emits open-loop action chunks and never checks the result, so I added a lower-rate visual supervisor over progress and terminal state that triggers stop, return, re-observation, or a bounded retry (terminal-verification F_1 **0.842**).
- Reached **70%** end-to-end success over 11 tasks (85% pick, 75% placement) with zero safety incidents; unseen object–target compositions are reported separately as an open condition, and built the voice-interaction and monitoring interface.

HONORS & AWARDS

James B. Angell Scholar, University of Michigan	Mar 2026
University Honors, University of Michigan (3 terms)	Dec 2024 – Dec 2025
Dean's List, University of Michigan (3 terms)	Dec 2024 – Dec 2025
Honorable Mention, Interdisciplinary Contest in Modeling (ICM)	May 2023
First Prize, 13th SJTU Liming Cup Mechanical Innovation Competition	May 2023
Outstanding Volunteer, Miyuan Volunteer Team	Mar 2023

SKILLS

Methods: trajectory optimization (DDP, augmented-Lagrangian methods), self-supervised representation learning (JEPA-style latent prediction), VLA fine-tuning & real-robot deployment ($\pi_{0.5}$), VLM integration

Languages: Python, C++, C, MATLAB

ML & robotics: PyTorch, Transformers, OpenCV, NumPy, SciPy, OpenAI Gym, ROS, PyBullet, Gazebo

Tools: Git, Linux (Ubuntu), Docker, $\text{\LaTeX}$, SolidWorks